

Makeline System

Hot-reloadable configuration for food assembly hardware

Quick Start

```
just generate-makeline simulation  
just simulate simulation
```

```
# Edit profiles/simulation.json  
# Ctrl+S → reloads in 2.5s
```

Aliases: `gm` / `s`

Always use `just generate-makeline` (not `just generate`)

Where Files Live

```
generated/profiles/simulation/  
├─ simulation.json          # Don't edit (gets regenerated)  
├─ spawner.json            # Process launch config  
└─ custom_layers.json      # Your customizations  
  
profiles/                  # Edit here  
├─ simulation.json         # Active profile (watched)  
└─ simulation/  
    ├─ makeline.json       # Expanded result  
    └─ backups/            # Timestamped copies
```

Edit `profiles/` **for hot-reload** | `spawner.json` launches processes

Profile Structure

Three sections define your makeline:

layouts: CabinetKind list (hardware topology)

- ```
{ "default": { "cabinets": ["Initial", "Denest", "Dispense",
"Lift"] }}
```

**layer\_groups:** Named edit collections (modifications)

- ```
{ "base": [layer1, layer2], "prod": [layer3] }
```
- Layers applied sequentially (order matters!)

line_builds: Combine layout + layer groups (final config)

Editing Workflow

Edit `profiles/simulation.json` → Save → 2.5s → Changed modules
reload

What happens: Diff → backup to `backups/{timestamp}/` → reload
affected modules

What keeps running: Planner, unchanged modules, active orders

Common Tasks

Change dispenser ingredient:

```
{ "EditSectionField": {  
  "identity": { "owner": "dispenser-3", "subject": "self" },  
  "section_name": "inputs",  
  "field_key": "assigned_ingredient_id",  
  "field_value": "black_beans"  
}}
```

Adjust buffer motion timeout:

```
{ "EditSectionField": {  
  "identity": { "owner": "buffer-1", "subject": "self" },  
  "section_name": "configuration",  
  "field_key": "motion_timeout",  
  "field_value": "10"
```

Physical Hardware Context

Real-world machine structure drives software design:

Cabinet = Physical enclosure unit

- Initial: System computer, no food hardware
- Denest: Unstacks bowls from dispenser
- Dispense: Contains ingredient hoppers (12-18 per cabinet)
- Lift: Presents finished bowls to customer

Device = Functional hardware subsystem within cabinet

- Core: Software-only (system services)

Architecture Rationale

Why this design?

Hardware abstraction: Physical machines have cabinets → devices → hardware modules. Software mirrors this hierarchy so config matches physical reality.

Hot-reload requirement: Must change config without stopping production. Graph-based design allows:

- Diff detection (compare old vs new graph)
- Selective module restart (only changed modules)
- Process isolation (spawner manages independent processes)

System Architecture

Three-tier enum hierarchy: CabinetKind → DeviceKind → ModuleKind

Generator expands each tier: Types define requirements → instances get created

CabinetKind enum (4 variants):

- Initial, Denest, Dispense, Lift

DeviceKind enum (10 variants):

- Core, CabinetCore, CabinetScreen, Denester, Dispenser

Cabinet → Device Mappings

Initial cabinet (system-wide services):

- `Core` device → 20 modules

Denest cabinet (bowl handling):

- `CabinetCore`, `Conveyance`, `Denester`, `CabinetScreen`

Dispense cabinet (ingredient dispensing):

- `CabinetCore`, `Hvac`, `DispenseFillPositioner`
- Plus $N \times$ `Dispenser` devices (12 for v5111, 18 for v5112)

Device → Module Details

Core device (Initial cabinet only) - 20 system modules:

- Api-1, BowlRecovery-1, CabinetMonitor-1, Datalog-1
- Discovery-1, Echo-1, Fault-1, Follower-1
- Interlock-1, LifeCycler-1, MachineConfig-1, Makeline-1
- PartnerApi-1, PartnerWebhook-1, Planner-1, Preprocessor-1
- RfidClient-1, Sequencer-1, State-1, Telemetry-1, Tracker-1

DispenseFillPositioner device (one per dispense cabinet):

- Buffer, Conveyance, Shutter, Duc

Core Modules Overview

System orchestration:

- **Planner**: Assembly planning (which dispenser, what order)
- **Sequencer**: Executes plans as bowl moves through system
- **Follower**: Tracks individual bowls (position, state)
- **Preprocessor**: Pre-processing incoming orders

Hardware control:

- **MakeLine**: Central coordinator, config hot-reload
- **MachineConfig**: Provides config to other modules

Graph Structure Deep-Dive

Graph representation: Directed acyclic graph (DAG)

Node types:

- `Root` (single)
- `Cabinet(CabinetKind)` (1-4 nodes)
- `Device(DeviceKind)` (variable count)
- `Module(module_data)` (40-100+ nodes)

Edges: Parent → child relationships

- Root → Cabinets

Generator Expansion Process

Step 1: Read `profiles/{preset}.json` → get layouts, layer_groups, line_builds

Step 2: Select line_build (from CLI or "default") → determines layout + which layer_groups to apply

Step 3: Expand layout into graph:

```
For each CabinetKind:
  cabinet_kind.devices() → Vec<DeviceKind>
  For each DeviceKind:
    device_kind.modules() → Vec<ModuleKind>
    Create numbered instances (buffer-1, buffer-2, ...)
```

Module Sections Reference

Different modules have different section names. Common sections:

Buffer modules: `configuration`

- Fields: `motion_timeout_ms`, `homing_velocity`, `home_to_lower_mrad`, etc.

Dispenser modules: `inputs`, `outputs`

- `inputs`: `assigned_ingredient_id`, `dispenser_kind`
- `outputs`: Runtime state (read-only)

Lifecycler module: `configuration`, `light`

Layer Edit Types - Part 1

EditSectionField: Change single config field (most common)

```
{ "EditSectionField": {  
  "identity": { "owner": "buffer-1", "subject": "self" },  
  "section_name": "configuration",  
  "field_key": "motion_timeout_ms",  
  "field_value": 20000  
}}
```

AssignSections: Replace entire sections (multiple related fields)

```
{ "AssignSections": {  
  "identity": { "owner": "lifecycle-1", "subject": "self" },  
  "sections": {  
    "configuration": { "cooldown complete ms": 24000, "timeout fault ms": 600000 },  
  }  
}
```


Layer Edit Types - Part 2

AssignDispensers: Populate all dispensers for cabinets

```
{ "AssignDispensers": {  
  "dispensers": [  
    { "cabinet_index": 2, "dispenser_index": 0, "ingredient_id": "black_beans",  
      "position": { "x": 0.0, "y": 0.0, "z": 0.0 }, "kind": "Auger" },  
    { "cabinet_index": 2, "dispenser_index": 1, "ingredient_id": "rice",  
      "position": { "x": 0.1, "y": 0.0, "z": 0.0 }, "kind": "Auger" }  
  ]  
}}
```

OmitModules: Remove modules from graph (testing/debugging)

```
{ "OmitModules": { "identities": [  
  { "owner": "hvac-1", "subject": "self" },  
  ...  
]}
```

Module Communication

IPC via ZeroMQ: Each module = separate process on pub/sub network

Identity for routing: Messages addressed by `{ owner, subject }`:

- `{ owner: "buffer-1", subject: "self" }` → message to buffer-1 process
- `{ owner: "buffer-1", subject: "motor-1" }` → message to buffer-1's motor child
- Makeline server routes based on Identity

Message types:

How Targeting Works

Identity structure: `{ owner: "module-name", subject: "target" }`

owner: Module instance name (numbered instances from graph expansion)

- `"buffer-1"`, `"buffer-2"` (dispense cabinet buffers)
- `"lifecycler-1"` (system singleton)
- `"lift-1"` (lift cabinet)
- `"dispenser-1"` through `"dispenser-12"` (v5111) or `"dispenser-18"` (v5112)

subject: Config target within the module hierarchy

Hot-Reload Mechanics

What triggers reload? Edit `profiles/{preset}.json` → `watch_server` detects → `makeline_server` reloads

Diff algorithm:

1. Load new profile, expand to graph
2. Compare new graph vs old graph (structure + sections)
3. Identify changed modules (section values differ)
4. Send `SectionChanged` events to affected modules
5. Modules reconfigure without restarting process

What causes module restart?

Custom Layers

Dev tweaks without modifying preset profiles

just generate-makeline-custom simulation → creates
custom_layers.json

```
[{
  "metadata": { "name": "Faster Buffer Motion" },
  "edits": [{
    "EditSectionField": {
      "identity": { "owner": "buffer-1", "subject": "self" },
      "section_name": "configuration",
      "field_key": "motion_timeout_ms",
      "field_value": 10000
    }
  ]
}]
```

Line Builds & Switching

One profile, multiple configs via different layer combinations

```
"line_builds": {  
  "production": { "layer_groups": ["base", "prod_ingredients"] },  
  "testing": { "layer_groups": ["base", "test_ingredients"] }  
}
```

At generation: just generate-makeline simulation testing

At runtime (Explorer):

- AvailableLineBuilds → see options
- SelectLineBuild { line_build_name: "production" } → switch (10- 22

Generated Files Explained

spawner.json: Process launch configuration

- Executable paths, args, environment variables for each module
- `generate-makeline`: Includes makeline_server references (enables hot-reload)
- `generate`: Includes machine_config references only (legacy, no hot-reload)

makeline.json: Expanded module graph

- Lives in `profiles/{preset}/makeline.json` (not generated!)
- Contains all module instances with their sections

- Adds hot-reload support via makeline_server in spawner.json

Watch mechanism: Only `profiles/{preset}.json` triggers 2.5s auto-reload

Spawner behavior: Smart restart - only kills/restarts processes with config changes

Never edit: `generated/` directory - gets overwritten on every generation

Available presets:

- `simulation` - Mock hardware, no real devices
- `v5111` - 4 cabinets, 12 dispensers

That's It

```
just generate-makeline simulation [line_build]
just simulate simulation
# Edit profiles/simulation.json
```

Key features:

- Hot-reload (2.5s, no restarts)
- Custom layers (dev tweaks)
- Line builds (test/prod switching)
- Auto backups (rollback ready)

Remember: `profiles/` not `generated/`